VE281 — Data Structures and Algorithms

Programming Assignment 1

Instructor: Hongyi Xin

— UM-SJTU-JI (Fall 2020)

1   Introduction

You are asked to implement six sorting algorithms: bubble sort, insertion sort, selection sort, merge sort, quick sort using extra array for partitioning, and quick sort with in-place partitioning.

You also need to implement an algorithm based on sorting algorithms: Graham's scan. It is a method of finding the convex hull of a finite set of points in the plane.

You will recall the knowledge about templates and function objects in this project, and also have a simple usage of STL (Standard Template Library).

A copy of the handout and the template of `<sort.hpp>` can be found in `/nfs/VE281-storage/P1` from the VM machines.

2   Programming Assignment

2.1   Sorting Algorithms

In `sort.hpp`, you need to write six functions with the six sorting algorithms:

- `bubble_sort`
- `insertion_sort`
- `selection_sort`
- `merge_sort`
- `quick_sort_extra` (quick sort using extra array for partitioning)
- `quick_sort_inplace` (quick sort with in-place partitioning)

You are allowed to add more helper functions in this file, but you're not allowed to change the definitions of these six functions defined by us.

Below is an example function `std_sort`, which use the `std::sort` function in the `<algorithm>` C++ standard library. However, do not copy and paste it into any other functions (you are prohibited from including additional libraries, such as `<algorithm>` in `<sort.hpp>`). Only use the below code in your main function to test your comparison function objects in your `<main.cpp>`.

```
template<typename T, typename Compare>
void std_sort(std::vector<T> &vector, Compare comp = std::less<T>()) {
    std::sort(vector.begin(), vector.end(), comp);
}
```

Recall what you've learned about STL containers and templates in VE280, here `T` is an arbitrary type which can be compared with the `comp` function (or *function object*). Each of your sorting function should provide the same result as the standard sorting function, which means for each pair of neighboring elements *a* and *b* in the sorted vector, a==b or `comp(a,b)==`true.

The `std::less<T>()` function is equivalent to the following function:

```
1   template<typename T>
2   bool compare_less(const T &lhs, const T &rhs) {
3       return lhs < rhs;
4   }
```

Alternatively, if `T` is restricted, i.e., you already know the first argument is `vector<int>`, you can define the compare function without using template:

```
1   bool compare_int_less(const int &lhs, const int &rhs) {
2       return lhs < rhs;
3   }
```

It can also be written in the form of *function object*:

```
1   template<class T>
2   struct CompareLess {
3       bool operator()(const T &lhs, const T &rhs) const {
4           return lhs < rhs;
5       }
6   };
```

Then you can call the sorting function with those methods (which are all equivalent):

```
1   std::vector<int> vector = {3, 2, 1};
2
3   std_sort(vector, std::less<T>());
4   std_sort(vector, compare_less<T>);
5   std_sort(vector, compare_int_less);
6   std_sort(vector, CompareLess<T>());
```

Note: `std::less` and `CompareLess` are *function objects*, so that they need to be initialized with `()`.

Important: do not include the compare function objects in `sort.hpp`. You are also not allowed to include additional libraries (in addition to the `vector` STL library) in `sort.hpp`. Remove all customized `include` statements from `sort.hpp` (including the ones that are commented out) before submitting to JOJ. Hint: include additional libraries in `main.cpp`.

Only submit the `sort.hpp` file (do not submit your main.cpp) to JOJ at `192.168.1.28:34765`.

2.2 Graham's Scan Algorithm[1]

2.2.1 Introduction

A set of points in a Euclidean space is defined to be *convex* if it contains the line segments connecting each pair of its points. The *convex hull* of a given set X may be defined as The (unique) minimal convex set containing X. For example, on a 2D plane, Figure 1 shows the convex hull of a set of points.


Figure 1: Convex hull of points on a 2D plane[2].

Graham's scan is a method of finding the convex hull of a finite set of points in the 2D plane. It is named after Ronald Graham, who published the original algorithm in 1972. The algorithm finds all vertices of the convex hull ordered along its boundary. It uses a stack to detect and remove concavities in the boundary efficiently.

In order to ensure the numeric precision, we use the sign of the cross product to determine relative angles when sorting and performing the algorithm. For three points $P_1(x_1, y_1)$, $P_2(x_2, y_2)$ and $P_3(x_3, y_3)$, compute the z-coordinate of the cross product of the two vectors $\overline{P_1 P_2}$ and $\overline{P_1 P_3}$, we define

$$ccw(P_1, P_2, P_3) = (x_2 - x_1)(y_3 - y_1) - (y_2 - y_1)(x_3 - x_1). \tag{1}$$

If the result is 0, the points are collinear; if it is positive, the three points constitute a "left turn" or counter-clockwise orientation, otherwise a "right turn" or clockwise orientation (for counter-clockwise numbered points).

2.2.2 Algorithm

The Graham's scan algorithm is defined in Algorithm 1.

Input : A finite set of points X

Output: A vector of points on the convex hull of X

$S \leftarrow$ empty stack;

$P_0 \leftarrow$ the point in X with the lowest y-coordinate (if the lowest y-coordinate exists in more than one point in X,
  the point with the lowest x-coordinate out of the candidates should be chosen);

sort the points in X by polar angle with P_0 (if several points have the same polar angle then only keep the
  farthest);

foreach P in X **do**

  while size(S) > 1 and `ccw`(next_to_top(S), top(S), P) $\leq$ 0 **do**

    pop from S;

  end

  push P to S;

end

<div align="center">Algorithm 1: Graham's scan.</div>

Now the stack contains the convex hull, where the points are oriented counter-clockwise and P_0 is the first point.

The polar angle of a point P with P_0 is the angle $\overline{P_0 P}$ forms with the x axis. In the sorting process, we compare the polar angle of two points P_1 and P_2 by computing $ccw(P_0, P_1, P_2)$. If it is positive, the polar angle of P_2 is greater than P_1. So the ccw function won't introduce any precision lost comparing to computing the cosine values of the angles.

In each iteration, we also use the ccw function, since we are popping the last point from the stack if we turn clockwise ($ccw < 0$) to reach this point in the while loop.

2.2.3 Input / Output

In `p1.cpp`, you need to implement this algorithm. You are allowed to use any of the sorting functions defined in `sort.hpp` (including `std::sort` from the `<algorithm>` library).

The input (from standard input) will contain $N + 1$ lines, where the first line is an integer N, and on the next N lines, there is a pair of point $P(x, y)$ on each line. All of the coordinates are integers in the range $[-2^{31}, 2^{31} - 1]$. The input would look like:

```
5
-2 -1
0 0
2 1
-3 4
5 -4
```

where each line contains the x and y coordinate, separated by a space.

The output (to standard output) should be the points outputted by the algorithm, which consist a convex hull of the input points. Each line should contain a point $P(x, y)$, which has the same format as the input. The first line should be P_0, and the order of these points are oriented counter-clockwise on the convex hull. This point sequence is unique, guaranteed by the algorithm. A sample output would look like:

```
5 -4
2 1
```

```
-3 4
-2 -1
```

You can assume that the number of points will always be correct.

2.2.4 Hints

You may test your sorting algorithms with the main program, but in your final submission, you are recommended to use `std_sort` for safe.

For the stack *S*, since you need to access the "next to top" element, you can use a `std::vector` instead of a `std::stack` for simplicity.

For the compare function, if you don't use a function object based one, global variables may be introduced, which is not a good idea. A better solution is to define a `struct` with an attribute P_0, and the `operator()` overload function.

When testing your algorithm, don't forget to include some extreme cases such as duplicated or collinear points.

2.3 Comparison of the Sorting Algorithms

We also ask you to study the performances of these six sorting algorithms, and also compare the results to the `std::sort` function.

For general *n* numbers, you should first generate different arrays of random integers with different sizes. Then you should apply your sorting algorithms to these arrays. Finally, you should plot a figure showing the runtime of each algorithm versus the array size *n*. For comparison purpose, you should plot all seven curves corresponding to the six sorting algorithms and the `std::sort` function in the same figure. (You do not need to upload the source codes for this comparison program.)

For sorting algorithms with $\mathcal{O}(n^2)$ time complexity such as bubble sort, when *n* is very large, it may take too long time waiting for the finish of the algorithm. You can set a threshold (i.e., 1 second) on the y-axis and skip the case.

For small number of input size ($n < 50$), run your algorithms many times in a single execution and compute the average time. Do you find anything interesting about the performance of your algorithms? Research and explain why `std::sort` is also very quick in small cases, while your quick sort function probably not.

Write a **short** (around 1 or 2 pages) report about your findings, plots and explanations.

2.3.1 Hints

- The performance of programs on Windows is usually not stable, so you should do the experiments on a Unix based system.

- You may want to write another program that calls the functions in `sort.hpp` to do this study.

- You can use the C++11 pseudo-random number generation library to generate more randomized numbers (instead of using `std::rand`).

- You can use the C++11 chrono library to get more accurate runtime of functions than `std::clock`.

- (Optional) You can use GNU Perf (only available on Linux) to find the bottleneck of your implementation.

- Although the major factor that affects the runtime is the size of the input array, however, the runtime for an algorithm may also weakly depend on the detailed input array. Thus, for each size, you should generate a number of arrays of that size and obtain the mean runtime on all these arrays. Also, for fair comparison, the same set of arrays should be applied to all the algorithms.

- You should try at least 5 representative sizes.

3   Implementation Requirements and Restrictions

3.1   Requirements

- You must make sure that your code compiles successfully on a Linux operating system with `g++` and the options `-std=c++1z -Wconversion -Wall -Werror -Wextra -pedantic`.

- You should not change the definitions of the functions in `sort.hpp`.

- You should only hand in two files `sort.hpp` and `p1.cpp`.

- You can use any header file defined in the `C++17` standard. You can use cppreference as a reference.

3.2   Memory Leak

You may not leak memory in any way. To help you see if you are leaking memory, you may wish to call valgrind, which can tell whether you have any memory leaks. (You need to install valgrind first if your system does not have this program.) The command to check memory leak is:

`valgrind --leak-check=full <COMMAND>`

You should replace `<COMMAND>` with the actual command you use to issue the program under testing. For example, if you want to check whether running program

`./main < input.txt`

causes memory leak, then `<COMMAND>` should be "`./main < input.txt`". Thus, the command will be

`valgrind --leak-check=full ./main < input.txt`

4 Grading

Your program will be graded along five criteria:

4.1 Functional Correctness

Functional Correctness is determined by running a variety of test cases against your program, checking your solution using our automatic testing program.

4.2 Implementation Constraints

We will grade Implementation Constraints to see if you have met all of the implementation requirements and restrictions. In this project, we will also check whether your program has memory leak. For those programs that behave correctly but have memory leaks, we will deduct some points.

4.3 General Style

General Style refers to the ease with which TAs can read and understand your program, and the cleanliness and elegance of your code. Part of your grade will also be determined by the performance of your algorithm.

4.4 Performance

We will test your program with some large test cases. If your program is not able to finish within a reasonable amount of time, you will lose the performance score for those test cases.

4.5 Report on the performance study

Finally, we will also read your report and grade it based on the quality of your performance study.

5 Acknowledgement

The programming assignment is co-authored by Yihao Liu, an alumni of JI and the chief architect of JOJ. This programming assignment is designed based on Weikang Qian's VE281 programming assignment 1.

References

[1] Graham scan - Wikipedia: https://en.wikipedia.org/wiki/Graham_scan

[2] Convex hull - Wikipedia: https://en.wikipedia.org/wiki/Convex_hull

Appendix

sort.hpp

```cpp
#ifndef VE281P1_SORT_HPP
#define VE281P1_SORT_HPP

#include <vector>

template<typename T, typename Compare>
void bubble_sort(std::vector<T> &vector, Compare comp = std::less<T>()) {
    // TODO: implement
}

template<typename T, typename Compare>
void insertion_sort(std::vector<T> &vector, Compare comp = std::less<T>()) {
    // TODO: implement
}

template<typename T, typename Compare>
void selection_sort(std::vector<T> &vector, Compare comp = std::less<T>()) {
    // TODO: implement
}

template<typename T, typename Compare>
void merge_sort(std::vector<T> &vector, Compare comp = std::less<T>()) {
    // TODO: implement
}

template<typename T, typename Compare>
void quick_sort_extra(std::vector<T> &vector, Compare comp = std::less<T>()) {
    // TODO: implement
}

template<typename T, typename Compare>
void quick_sort_inplace(std::vector<T> &vector, Compare comp = std::less<T>()) {
    // TODO: implement
}

#endif //VE281P1_SORT_HPP
```